

Noisy Language Reference, Draft Version 0.2

Phillip Stanley-Marbell
psm@mit.edu

MIT CSAIL,
Cambridge, MA 02139.

Sensor-driven hardware platforms require appropriate programming abstractions. Noisy is a language designed for programming sensor-driven hardware platforms such as the MIT Warp platform. This document describes the motivation and design of the language and provides examples of its use.

Contents

1	Introduction	2
1.1	Signals and physics in Noisy	4
1.2	Dimensions and units of measure in Noisy	5
1.3	Time series of signals in Noisy	5
1.4	Notating and computing with timing uncertainty and measurement value uncertainty in Noisy	5
1.5	Organization of Noisy applications	7
2	Lexical Elements	7
2.1	Reserved tokens	7
3	Syntactic Elements	8
3.1	Programs	9
3.2	Types, type expressions, and type declarations	9
3.3	Variable and Channel Identifiers	11
3.4	Channel declarations versus channel definitions	11
3.5	Structuring programs	11
3.6	Scopes	13
3.7	Statements	13
3.8	Assignment statements	14
3.9	Expressions	14
3.10	The <code>match</code> and <code>matchseq</code> constructs	14
3.11	The <code>iter</code> construct	16
4	Operator Descriptions	17
4.1	Operators on names	17
4.2	Operators on channels	18
5	Types	18
5.1	Reference versus value types, channels and <code>nil</code>	18
5.2	Specifying tolerance constraints	18
5.3	Types <code>bool</code> , <code>nybble</code> , <code>byte</code> , <code>int</code> , <code>real</code> and <code>fixed</code>	20
5.4	Type <code>string</code>	20
5.5	The <code>array</code> collection type	20
5.6	Tuple collection type	21

5.7	Aggregate Data Type (ADT) collection type	21
5.8	The <code>set</code> collection type	21
5.9	Recursive list collection type	22
5.10	Dynamic data structures	22
5.11	The <code>vector</code> type	23
6	The Runtime System	23
6.1	Runtime system abstraction	23
6.2	Target hardware model	23
6.3	Structure of the runtime name space	24
6.4	Enforcing access restrictions on names	24
6.5	Runtime system implementation	24
	Appendix	27
A	A call-by-value Lambda calculus with error-tolerant types	27
B	Formal Definition of Type System	28
C	Implementation	28
C.1	Building the AST	28
C.2	Modifying the AST for easier code generation	29
C.3	The symbol table	29
C.4	Types	29
D	Noisy Language Grammar	30
	References	32

1. INTRODUCTION

Platforms driven by sensors are an increasingly important computing domain. These platforms range from wearable health-tracking devices, to home or office monitoring devices, to consumer and commercial robots, to unmanned aerial vehicles or drones. These platforms, which are referred to variously as either *cyberphysical* systems, *embedded* systems, or *Internet-of-Things* platforms, have a common property: Their software processes data from physical world interfaces such as temperature sensors, accelerometers, humidity sensors, or global positioning system (GPS) modules, and they may modify their environments through actuators such as motors or relays, or interact with users through displays.

Today, in order to be able to interact hardware, which typically occurs through memory-mapped I/O to microcontroller peripherals such as I2C [24] or SPI [23], most of the software for these platforms is written in a combination of C and assembly language. Higher-level control and logic for these platforms is typically also written in C or C++, or may be written in languages such as Python. To help bridge the gap between facilities provided in traditional programming languages and runtime systems, and the specific constraints of sensor-driven systems, several commercial and research language designs and implementations have explored limited support for sensor-driven systems, such as adding units of measure or fundamental and derived dimensional quantities to augment program variable types and to check these at compile and runtime [17; 7; 3; 10; 4; 25; 20; 32; 11; 14; 15; 5; 18; 19].

Today, no existing programming languages, libraries, or runtime systems for sensor-driven systems provide language constructs or standard libraries to purposefully exploit restrictions caused by the laws of physics on the properties of the signals they process. Existing languages provide no mechanisms to ease the implementation of important techniques such as sensor fusion for combining the data from multiple sensors or

exploiting correlations between sensors mounted in a known physical configuration to obtain better accuracy, lower noise, or lower energy usage. And, existing languages and runtime systems do not provide constructs to assist programmers in modeling important concepts such as measurement uncertainty or significant figures. There is therefore an unmet need for software libraries, new programming constructs, or new programming languages where appropriate, to allow programmers to more easily deal with measurement data from the physical world, which is inherently noisy.

Noisy is a language for processing signals from the real world, and for exploiting knowledge about correlations between signals or invariants obeyed by signals, to make the task of programming easier, to enable new types of compile-time program transformations based on physics, and to make programs more efficient. The design motivations for Noisy are:

- (1) **Coupling signals to program variables:** To allow programmers to denote the physical signals associated with program variables.
- (2) **Coupling signals to invariants and physical laws:** To allow the compilers and runtime system to exploit coupling between signals through invariants (laws) relating signals specified in the Newton [22] language.
- (3) **Basic data types for signal processing:** To ease the work of programmers implementing computations on values from sensors by providing native datatypes for time series, vectors, sets, and operators that ease the computation on these physical types such as computing the frequency modes of time series data, integrating, and differentiating time series data.
- (4) **Measurement uncertainty:** To allow programs to determine the uncertainty of values of variables representing signals as they are read directly from sensors or after subsequent computing on them.
- (5) **Tolerable result uncertainty:** To allow programs to specify precision, accuracy, reliability, significant figures, latency tolerance constraints, and constraint violation control flow.
- (6) **Program transformations exploiting signal correlations, physics, and uncertainty:** To incorporate an appropriate set of facilities into the language to enable compilers to extract semantic information necessary to enable new program transformations relevant to sensor-driven systems.

Noisy is intended to be an implementation platform to investigate a number of research questions. The specific research aims include:

- **Substituting sensors in computations based on physical correlations:** To evaluate whether sensor substitution is a viable technique for either improving the accuracy of sensing algorithms or for reducing the energy usage of sensor driven programs while maintaining their functionality.
- **Automating opportunities for reducing sensor accuracy:** To evaluate whether tracking measurement uncertainty can provide automated opportunities for purposefully inducing measurement uncertainty for sensor interactions in exchange for lower sensor power dissipation, and to quantify how much benefit can be gained from doing so for realistic programs.
- **Safety of sensor-driven programs:** To evaluate whether tracking measurement uncertainty in programs can make programs safer. Our goal is to evaluate the effect of significant figure tracking on program safety both by quantitatively evaluating the effect of compiler feedback enabled by language-level significant figure tracking, as well as safety enabled by program transformations and compile-time analysis guided by language-level significant figure annotation. One measure of program safety will be whether our proposed techniques allow us to identify or repair errors in example programs known to have defects.
- **Implementation efficiency of significant-figure tracking:** To evaluate whether tracking measurement uncertainty in programs can be achieved efficiently.
- **Utility and overhead of signal processing types:** To evaluate the utility of language-level support for signal processing such as representing time series and operators on time series, and to quantify the performance benefit that these facilities can provide by exploiting hardware opportunities for signal processing acceleration.

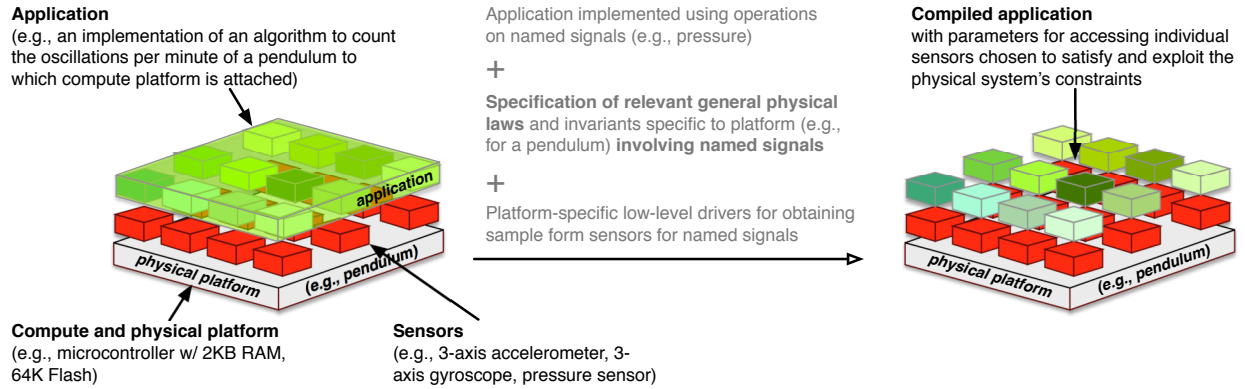


Fig. 1. Programs written in Noisy specify interaction with sensors in terms of signal types such as **pressure** or **temperature**. Hardware system designers provide physical platform descriptions written in Newton which specify general physical laws relevant to the physical system as well as specific invariants satisfied by a given physical system. Along with low-level device interface code for retrieving samples from sensors, the Noisy program is compiled to executables for a given platform. The transformations applied to the Noisy program are enabled by the information exposed in the Newton platform description.

Figure 1 shows how Noisy programs are combined with a Newton platform description and low-level device interface code and compiled to executables for a given platform.

```
1 p1 : int pressure;
```

Fig. 2. Numeric types in Noisy can include an optional signal designation.

1.1 Signals and physics in Noisy

Noisy introduces the concept of *signal designations*. Any numeric type in Noisy can be designated to be one of seven base dimensions corresponding to the seven base SI units, one of five base signals (derived dimensions) (**pressure**, **acceleration**, **magneticflux**, **relhumidity**, **anglerate**), or one of the additional signals defined in a platform's Newton description, described further in later sections. Figure 2 shows an example of a declaration of a numeric variable of type **int** with a signal designation of **pressure**. Noisy's semantics requires variables with signal designations to be used in *dimensionally-consistent* ways: addition, subtraction and comparisons of variables with different signal designations are not well-typed.

Noisy's semantics treats numeric types with signal designations as obeying any invariants specified in a platform's Newton description. This allows a Noisy compiler to use the invariants specified in a platform's Newton description perform program transformations that satisfy those invariants but which may improve program performance or efficiency.

To read sensor values for a given signal type, programs declare *channels* with the type of the desired signal, and read from those channels. Channels in Noisy are described later in this document. Reads from signal channels not supported on a given platform fail, returning the value **nil**, in much the same way that memory allocation in languages with dynamic memory allocation will fail if the platform does not have sufficient memory.

```
1 h : int, dimensions distance;
2 p2 : int, signal pressure, units (1/1000) kilogram meter**(-1) second**(-2);
3 seaLevelPressure : const 101325, units kilogram meter**(-1) second**(-2);
```

Fig. 3. Numeric types in Noisy can include an optional signal designation.

1.2 Dimensions and units of measure in Noisy

Variables with signal designations implicitly acquire the signal's SI units of measure. For example, the variable **p1** in Figure 2 has implicit units $\text{kilogram} \cdot \text{meter}^{-1} \cdot \text{second}^{-2}$ or **kilogram meter**(-1) second**(-2)**, where the operator ****** denotes exponentiation. Figure 3 shows how variables and constants may also be assigned explicit dimensions and units of measure. For example, in Figure 3, the variable **h** is declared with explicit dimensions of length, and the variables **p2** and **seaLevelPressure** both have the same dimensions ($\text{mass} \cdot \text{length}^{-1} \cdot \text{time}^{-2}$) but have different units or measure.

Values without explicit or implicit units of measure are unitless. Noisy's semantics requires all expressions in comparisons, addition, subtraction and assignment to be of the same units of measure. Under these semantics, the variable **p2** and the variable **seaLevelPressure** have different units of measure (units **(1/1000) kilogram meter**(-1) second**(-2)** or milli Pascals for **p2**, in contrast to units **kilogram meter**(-1) second**(-2)** or Pascals for **seaLevelPressure**). Noisy provides operators for evaluating the units of measure of variables and constants and these are described further in later sections.

```
1 f : (integerArgument: int, dimensionParameter: dimparam) : (resultValue: dimensionParameter**2)
2 {
3 }
```

Fig. 4. Numeric types in Noisy can include an optional signal designation.

Subroutines in Noisy may be parametrized on dimensions. Figure 12 shows a subroutine, **f**, which takes a value **integerArgument** as argument and a dimension parameter **dimensionParameter** and returns a result **resultValue** with dimensions **dimensionParameter**2**.

```
1 pressureSamples : int, signal pressure, timeseries;
2 gyroSamples    : int, signal anglerate, timeseries;
3
4 gyroSamples    <- [kSampleCount] signal anglerate;
5 pressureSamples <- [kSampleCount] signal pressure;
```

Fig. 5. Time series in Noisy simplify processing sensor signals.

1.3 Time series of signals in Noisy

Values read from sensors are often time series of readings sampled at a known rate. Noisy allows programmers to define variables as time series using the **timeseries** designator. Reads from channels can specify a number of samples to read into a time series variable. Figure 5 illustrates with an example. Noisy provides several operators for operating on time series data, such as for obtaining the list of frequency and amplitude tuples of the modes of oscillation in a signal (**periods**), integration over time and differentiation with respect to time (**tintegrate** and **tdifferentiate**), and obtaining a list of the time stamps or samples for a given time series (**timebase** and **samples**) as well as constructors for creating time series from a list or array of numeric values (**timeseries**).

1.4 Notating and computing with timing uncertainty and measurement value uncertainty in Noisy

All physical measurements have limited precision and as a result, all signal sample values have a limited number of significant figures. As values are used in the arithmetic operations of signal processing algorithms, and combined through arithmetic with constants and with other signal sample values, the uncertainty in the individual bits of their value representations may degrade further. Existing languages provide no way to determine the amount of measurement uncertainty associated with values obtained from sensors. Recent work [9] provides one mechanism for tracking measurement uncertainty once values are obtained from sensors by keeping track of the entire probability distribution of values and using Bayesian networks constructed based on a program's dataflow, to compute new probability distributions for each step of a computation.

Such methods of algebra on arbitrarily-distributed random variables is seductive, but is only possible for a limited number of types of distributions and operations.

```

1 pressureSample      : int, signal pressure;
2 seaLevelPressure    : const 101325, units kilogram meter**(-1) second**(-2);
3 approximatePressure : int, signal pressure, sigfigs 10;
4
5 gyroSample          <- signal anglerate;
6
7 /*
8  *      The <> operator returns the number of significant figures
9  *      of its operand.
10 */
11 if (<gyroSample> == <seaLevelPressure>)
12 {
13 }
14 else if (<gyroSample> == <approximatePressure>)
15 {
16 }

```

Fig. 6. By allowing programmers to specify the number of significant figures to be maintained for variables, and by providing constructs to determine the number of significant figures in signal samples obtained from sensors, Noisy makes it easier for programmers to write safer software that processes signals from the real world.

Noisy provides programmer support for reasoning about value uncertainty in sample values of signals as well as any numeric program value using three techniques: Static source-level annotation of significant figures for constant and variables, a dynamic runtime semantics which tracks the number of significant figures for every program value, and dynamic runtime determination of significant figures for any value in a Noisy program. The Noisy runtime system implements this semantics by maintaining at most one an additional memory word for each live numeric variable in a Noisy program. In practice, the overhead can be reduced by when the Noisy compiler can statically determine that information on the number of significant figures of a given variable is never used at runtime. Figure 6 illustrates how Noisy programs can specify the significant figures in a constant, or can determine the number of significant figures in a value obtained at runtime from a sensor signal. The syntactic constructs for source-level annotation of significant figures and the semantics of Noisy’s significant figure constructs are described later in this document.

In addition to specifying the number of significant figures that are meaningful for a given program variable, programmers may wish to specify that values of specific program variables may take on a given distribution of random errors at runtime.

We refer to the difference between the correct and erroneous values as the *error magnitude*. An error magnitude tolerance specifies the acceptable distribution on the error magnitude values. For example, a plain English specification of a simple error tolerance constraint could be “Ensure that the error magnitude only exceeds 2.0 at most one out of every million times that a value is read”. If the error magnitude is denoted by a random variable M , the above statement could be written more precisely as

$$\Pr\{M > 2.0\} \leq \frac{1}{1,000,000}$$

When channels in Noisy act as the interface to sensors, reads from a channel may fail due to communication failures with a sensor [30]. Such loss will be referred to as *channel-level erasures*. A channel erasure tolerance constraint thus specifies the acceptable distribution on the *loss fraction* of values.

Because operations on channels, can have non-deterministic latencies, since they depend on the state of the underlying communication interconnect, a *latency tolerance constraint* enables a Noisy program to specify how much latency on a communication is acceptable. The syntactic notation for these error-tolerance specifications are described later in this document.

Error tolerance specifications allow the Noisy compiler to perform two kinds of transformations. First, given an error-tolerance specification, the Noisy compiler can statically at compile time deduce the implied number of significant figures, and this can be queried by programs dynamically at runtime as described above. Second, given an error-tolerance specification on a channel variable, the Noisy compiler can use allow samples to be elided at runtime, or sample values to be encoded using techniques such as value-deviation-bounded serial (VDBS) encoding [31; 28] to reduce sensor I/O energy costs. Noisy’s error-tolerance specification constructs are described later in this document.

1.5 Organization of Noisy applications

Programs in Noisy are organized into units called *name generators* (*namegens*) (Figure 7). Name generators are collections of program statements (e.g., like functions or procedures in Algol family languages) which exchange information by explicit communication rather than transfer of control flow. They interact by communicating on *names* by which each name generator is represented in a runtime *name space*. Namegens and names are analogous to *Actors* and their *mail addresses* in the Actor system [1]. All interactions on names are expressed in terms of a small alphabet of operations, forming the *name interaction protocol*. Names abstract entries in the name space, which abstract channels, in much the same way that pointers abstract memory addresses, which abstract memory cells. Decoupling the interactions between namegens through entries in the runtime name space, facilitates the easy migration of endpoints of communications. This is beneficial to application re-mapping in the presence of detected failures, for a programmable substrate with redundantly available resources.

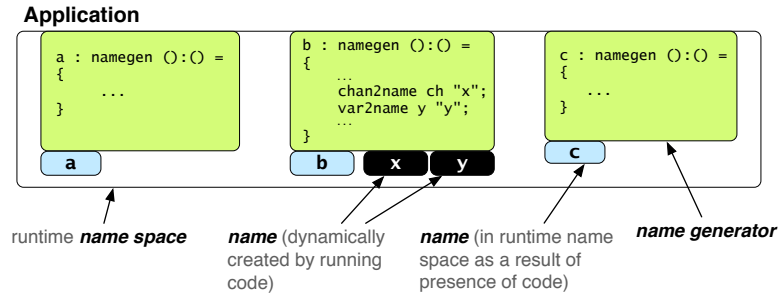


Fig. 7. Organization of applications into *namegens* which interact via *names* in a *runtime name space*, using a simple *name protocol*.

2. LEXICAL ELEMENTS

Programs are sequences of Unicode [33] characters, organized according to the rules of the language grammar. Tokens are separated by whitespace, operators or other separators. Whitespace consists of '_ ', '\n', '\r' and '\t'. Separators and operators may abut one or more tokens. In addition to these separators and operators, several tokens are reserved for use in the language and have special meaning; they may not be used as identifiers. Comments are introduced with the character #, and continue until the next newline, or the end of input:

```
#      This is a comment
```

2.1 Reserved tokens

A set of single and multiple character tokens are lexically reserved and may not be used in identifiers. The reserved operators and separators are:

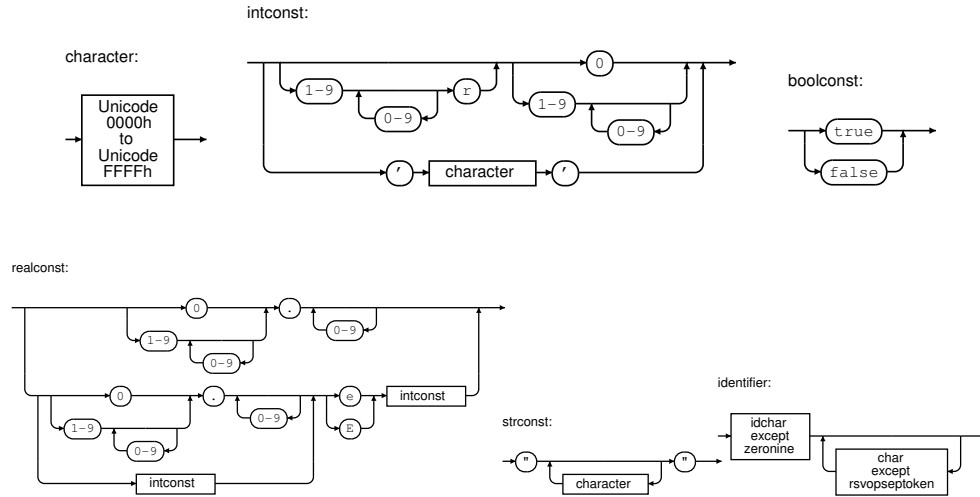


Fig. 8. Syntax diagrams for lexical elements.

~	!	%	^	&	*	()	-	+
=	/	>	<	;	:	,	"	{	}
[]		<-	.	,	<=	>=	^=	=
&=	%=	/=	*=	-=	+=	:=	!=	>>	>>=
<<	<<=	&&		::	==	--	++	<-=	=>
->	#	==@	0<	**					

The reserved identifiers in the language are:

acceleration	adt	alpha	ampere	anglerate	bandpass	bandstop
bool	byte	candela	chan	chan2name	const	crossproduct
current	dimensions	dimparam	distance	dotproduct	else	epsilon
erasures	errors	false	fixed	fn	for	hd
highpass	if	int	iter	kelvin	kilogram	latency
len	list	lowpass	luminosity	magneticflux	mass	match
matchseq	material	meter	mole	name2chan	namegen	nil
nybble	of	periods	predicate	pressure	proctype	rat
real	relhumidity	reverse	samples	second	set	sigfig
signal	sort	string	tau	tderivative	temperature	time
timebase	timeseries	tintegral	tl	true	type	units
universe	valfn	var2name	vector			

The production rules employed in lexical analysis for determining the terminal symbols from sequences of Unicode characters, can be described by a *regular language* [2], whose productions are shown by the syntax diagrams in Figure 8.

3. SYNTACTIC ELEMENTS

A sequence of Unicode characters, which corresponds to a sequence of one or more tokens as described in the preceding section, is a valid program if it conforms to the rules of the language grammar. A complete language grammar in EBNF [35] notation is provided in Appendix D.

3.1 Programs

A Noisy program is composed of a single program interface definition (*proctype definition*) and a collection of *name generators* (*namegens*). The proctype definition specifies a set of constants, type declarations, and publicly visible namegens. All entries in the proctype definition become visible in the runtime names space.




Fig. 9. Syntax diagram for overall program structure.

Example:

```

1 Hello : proctype
2 {
3     init : namegen();
4 };

```

3.2 Types, type expressions, and type declarations

Noisy is statically checked and strongly typed. Every expression can be assigned a unique type, and the compiler will not accept as valid any program for which a type error can occur at runtime. The syntax of the basic and derived types are represented by the syntax diagrams in Figure 10. The basic data types are `bool`, `nybble`, `byte`, `string`, `int`, `real`, and `fixed`. A basic type may have associated with it one or more *error*, *loss* (i.e., erasure), or *latency* tolerance constraints. Variables of type `bool` are 1-bit values, `nybble` are 4-bits, `int` are 32-bit signed integers in two's complement format. Type `real` is a 64-bit double-precision floating point value in IEEE-754 format. The type `fixed` is a 16-bit fixed-point representation for real values.

Types can have *dimension designations* corresponding to one of the seven base SI units (**distance**, **mass**, **time**, **material**, **current**, **luminosity**, **temperature**).

Types can have *units of measure designations* corresponding to one of the standard units for the seven base SI units (**meter**, **kilogram**, **second**, **mole**, **ampere**, **candela**, **kelvin**).

Types can have *signal designations*. Any numeric type can be designated by a program to be one of seven base dimensions corresponding to the seven base SI units (**distance**, **mass**, **time**, **material**, **current**, **luminosity**, **temperature**), one of five base signals (derived dimensions) (**pressure**, **acceleration**, **magneticflux**, **relhumidity**, **anglerate**), or one of the additional signals defined in a platform's Newton description.

The type collections are arrays, abstract data types (ADTs), lists, tuples and sets. Arrays are sequences of elements of a single type. ADTs and tuples are collections of elements of possibly different types; whereas an ADT collection has an associated *type name*, tuples are unnamed collections. Lists are collections of elements of a single type, with access to only the head of the list. Sets are unordered collections of elements, with primitive operations for union, intersection, relative complement and cardinality, with which idioms for other set operations (e.g., membership, subset) can be implemented¹. The language types are discussed in more detail in Section 5, and a formal description of the type system is provided in Section B.

The types `bool` and `nybble` are motivated by the fact that there are many classes of applications in which single bit values are semantically popular, e.g., for use as flags, as are variables taking on a small range of values [27]. In memory constrained devices, it makes sense to provide language constructs that support these idioms. Particularly for small cardinality, sets provide a means for memory and computationally efficient implementation of collections. Example:

¹Wirth [36] attributes the idea of sets as a language data type to Hoare, circa 1972.

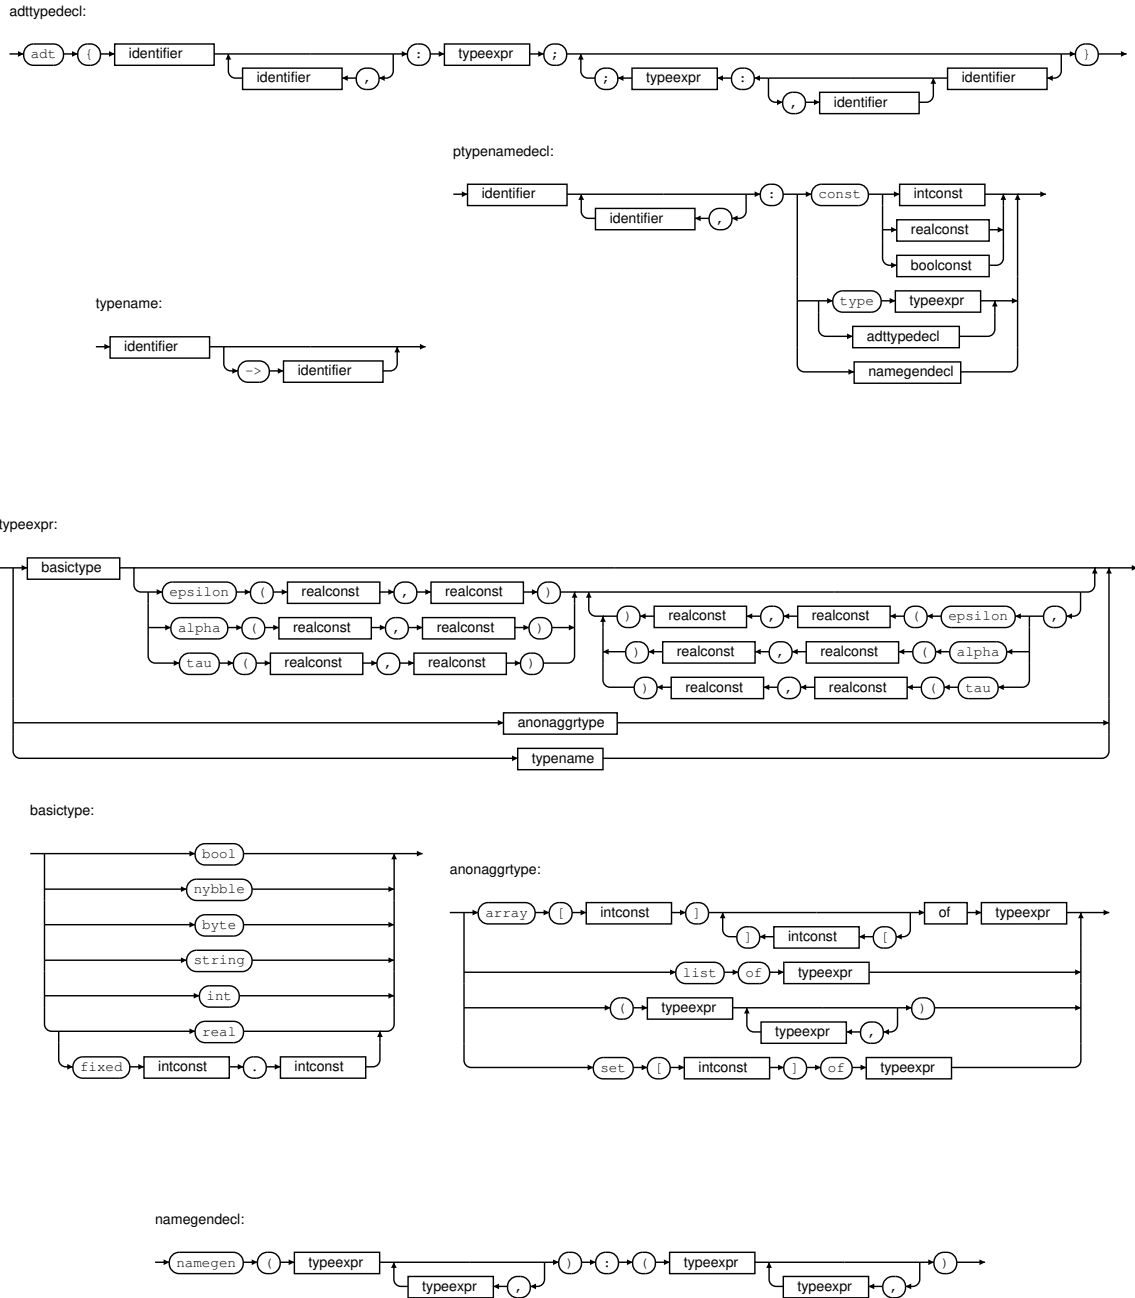


Fig. 10. Syntax diagrams for types.

```

1 # Variable 'pixel' is an integer with constraint that the
2 # probability that the magnitude of error in its value being > 1
3 # is less than 0.01
4
5 pixel : int, epsilon(1, 0.01);

```

3.3 Variable and Channel Identifiers

There are two kinds of program-defined identifiers in Noisy: *variables* and *channels*. Variables are identifiers that are used to refer to possibly structured data in memory. The simplest type of variable is a *pronumeral* that represents an item in memory with one of the basic arithmetic (numeric) types. Channels on the other hand are identifiers that are used to refer to communication paths between namegens. Channels as a programming language construct trace their roots to Hoare's *Communicating Sequential Processes (CSP)* [16]. Noisy channels are similar to channels in Alef [], Limbo [], Go [] and to continuation variables in Cilk [8]. Noisy channels are tied to a *name* in the runtime name space.

While variables can be used in any expression (Section 3.9), channels can only be used in *send* and *receive* expressions. In a send expression, a value is written to a channel, while a value is read from a channel in a receive expression. The evaluation of a send expression does not complete until *another* namegen performs a receive operation on a channel that is associated with the same name in the runtime name space as the channel being written to. Likewise, evaluation of a receive expression will not complete until a send expression on a channel associated with the same name in the runtime name space is executed by another namegen.

Variables and channels must be defined before use, ascribing to the variable or channel a basic or collection type. A variable may also be declared in conjunction with assignment, in which case it is ascribed the type of the value it is being set to. Example:

```

1 #   Variable 'sensor' is a channel of real valued numbers,
2 #   with the constraint that the probability that the number
3 #   of lost values on the channel being > 1 is 0.1, and > 10 is 0.0003
4
5 sensor : chan of int, tau(1, 0.1), tau(10, 0.003);
6
7
8 #   Declare variable 'a' with type int
9 a      : int;
10
11 #   Declare variable 'b' with tuple type (int, int)
12 b      : (int, int);
13
14 #   Define variable 'c' with type being the type of its assigned value (real)
15 c      := 1.0;

```

namegendefn:




Fig. 11. Syntax diagram for name generator definition.

3.4 Channel declarations versus channel definitions

Channels are used to communicate between two namegens. Before a channel can be used, it must be associated with a name in the runtime name space via a *name2chan* expression. A channel declaration only serves to ascribe a type to the identifier that will eventually refer to a name in the runtime name space. A channel *definition* is the association of a name in the runtime to a channel identifier, i.e., assigning the result of a *name2chan* to an identifier — prior to a definition, channel identifiers have the value *nil*.

3.5 Structuring programs

Programs are composed of *program types (progtypes)* and their implementations, a collection of *name generators (namegens)*. A progtype is a unit of modularity that defines a set of types, constants, and name generators. The

unit of compilation of programs is a single progtype and its implementation. This single complete program input to the compiler is used to generate one or more compiled outputs, corresponding to the pieces of the partitioned application. Partitioning occurs most easily at the level of individual namegens. Due to the structure of the language and its runtime system, it is however possible to further partition a single namegen into smaller pieces.

A *name generator* or *namegen* is a collection of type declarations, variable definitions and statements. It is the unit at which applications get partitioned into executable units. There is no shared state between namegens. Namegens are not functions: i.e., control flow does not pass between namegens. A namegen may bind a channel variable to a *name*. Syntactically, names are represented by strings. Such strings are part of a runtime *name space*². The syntax diagrams for the grammar production for a valid namegen definition is shown in Figure 11.

All interaction between namegens is explicitly through *names* in the runtime name space, via the binding of names to channels. A namegen definition includes an interface tuple type. This is the type structure of the name for write and read operations on the name. Example:

```

1 Math : progtype
2 {
3     sqrt : namegen (real) : (real);
4     exp  : namegen (real, real) : (real);
5 }
6
7 # The actual implementation of the Math progtype must have
8 # definitions for the namegen types sqrt and exp:
9 sqrt =
10 {
11 }
12
13 exp =
14 {
15 }
16
17 # A namegen that does not appear in the interface is declared
18 # and implemented in one step here
19 somefunc : (real, real) : (real) =
20 {
21     # Declaration of 'sqrt' in progtype introduced a new type name.
22     # We use that as the type of the variable bound to the namegen sqrt:
23     s := name2chan sqrt "sqrt";
24 }

```

Subroutines in Noisy may be parametrized on dimensions. Figure 12 shows a subroutine, **f**, which takes a value **integerArgument** as argument and a dimension parameter **dimensionParameter** and returns a result **resultValue** with dimensions **dimensionParameter**2**.

```

1 f : (integerArgument: int, dimensionParameter: dimparam) : (resultValue: dimensionParameter**2)
2 {
3 }

```

Fig. 12. Numeric types in Noisy can include an optional signal designation.

The definition of a namegen in a progtype introduces a new type name made up of a *write type* and a *read type*, that can be used in a name2chan expression. If a namegen is not defined in its progtype declaration, it is not visible within the name space once the progtype implementation is loaded. However, the definition of the namegen within the body of the program introduces a new type, as it would if it had been declared in the progtype. Within the body of a namegen, the namegen's name is a channel which when read from has the namegens *read type*, and which

²By analogy, a computing system fashioned on the von Neumann model has the concept of *memory addresses* which are represented by integers, and are part of an *address space*.

for the purposes of writes, has the namegen's *write type*. Applying the `name2chan` operator to a name that has type `namegen` causes the instantiation of a new thread of control / instance of the namegen. Example:

```

1
2 #       Defines a namegen mul.  An antry in the name space "mul"
3 #       with write type (int, int) and read type (int) will exist
4 #       at runtime
5
6 mul : (int, int) : (int) =
7 {
8 }

```

3.6 Scopes

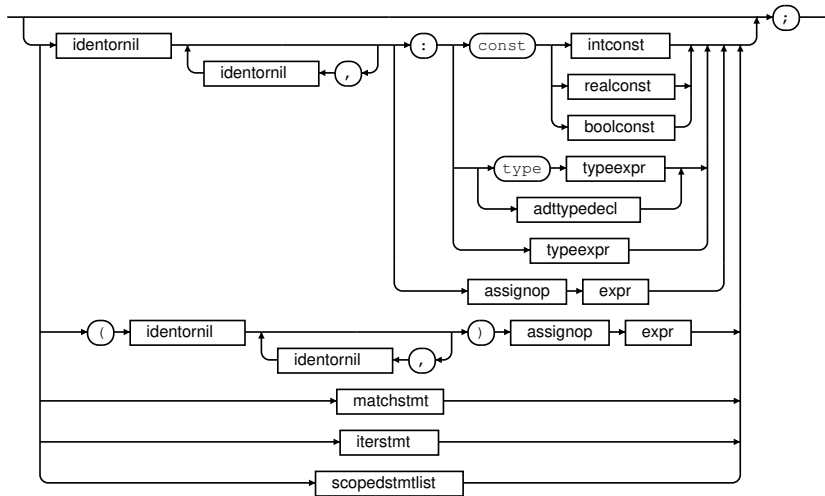
The largest scope is the body of a namegen; there is no global scope. Within a namegen, a new scope is embodied by a collection of statements enclosed in braces (`{ ... }`). Variables defined within a scope have visibility only within the given scope. The following grammar productions introduce new scopes:

- `proctypebody` production
- body of ADT type declaration
- `scopedstmtlist` production
- `guardedscopelist` production


The complete language grammar is provided in Appendix D.

3.7 Statements


stmt:



identornil:



scopedstmtlist:



assignop:




Fig. 13. Syntax diagrams for statements

The statements which make up the body of a namegen define how it computes, creates names in the runtime name space, and communicates with other namegens. The syntax diagrams for grammar productions corresponding to valid statements is shown in Figure 13.

3.8 Assignment statements

The *r-value* in an assignment must be of the same type as the *l-value*, with the exception of an *l-value* of nil, which can be assigned an expression of any type; the *r-value* in such an assignment to nil is first evaluated, but the assignment operation does not yield any further action. The basic assignment operator is =. The additional assignment operators which have the form *binop*= assign to the *l-value* the result of *l-value binop r-value*. The syntax diagrams for the associated grammar productions are illustrated in Figure 13. Example:

```

1
2 #      Bit-wise negation of variable 'a'
3 a ^= 1;
```

3.9 Expressions

The associativity of operators is implicit in the definition of the grammar for expressions: *expressions* are made up of *terms* and the low precedence binary operators; terms are made up of *factors* and the high precedence binary operators; factors are *l-values*, constants, expressions in parenthesis or unary operators followed by a factor. The unary operators thus have the highest precedence, followed by high precedence binary operators and then low precedence binary operators. This is illustrated by the grammar productions in Figure 14. Operators are discussed in more detail in Section 4.

3.10 The match and matchseq constructs

The **match** statement is a collection of guarded statement blocks. It executes *all* constituent statement blocks whose guards, which are Boolean expressions, evaluate to true. If multiple guards evaluate to true, the order in which the guarded statement blocks are evaluated is non-deterministic. The **match** statement is analogous to *guarded selection* in Dijkstra's guarded commands [12]. The **matchseq** statement on the other hand evaluates its guards sequentially, until a guard that evaluates to true. The guarded statements are then executed, and the **matchseq** statement completes. The syntax of **match** and **matchseq** statements are illustrated by the syntax diagrams in Figure 15.

The **match** statement can be used to implement the equivalent of if statements (as in the Pascal and C family of languages), and multi-way selection statements such as Pascal **case** and C **switch** statements. Example:

```

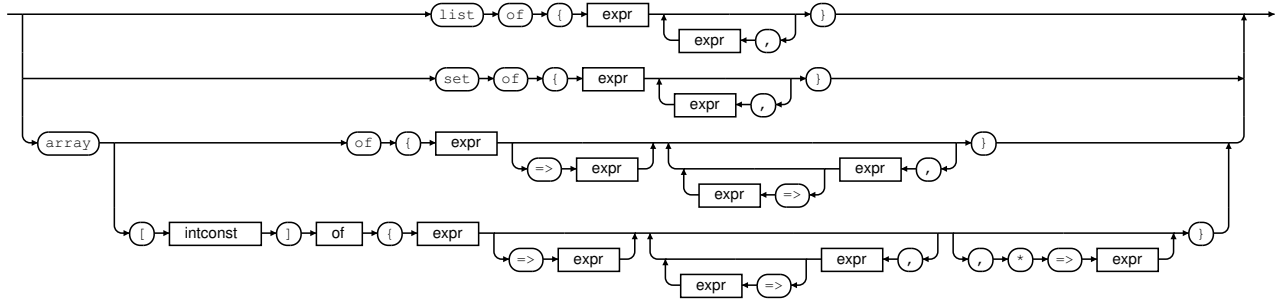
1
2 #      If 'a' > 'b', 'max' = 'a',
3 matchseq
4 {
5     a > b  =>    max = a;
6     true   =>    max = b;
7 }
```

One possible implementation of a Fibonacci sequence generator namegen:


```

1
2 fibonacci : progtype
3 {
4     fib : namegen (int):(int);
5 }
6
7 fib : (int) : (int) =
8 {
9     v := <-fib;
10
11     matchseq
12     {
```


anonaggrcastexpr:



expr:




factor:




chanvtxexpr:



term:




chan2nameexpr:



var2nameexpr:



fieldselect:



name2chanexpr:




Fig. 14. Syntax diagrams for expressions

```

13      (v == 0) =>
14      {
15          fib <-= 0;
16      }
17
18      (v == 1) =>
19      {
20          fib <-= 1;
21      }
22
23      true =>
24      {
25          c1, c2 := name2chan int "fib" 0.0;
26          c1 <-= (v - 1);
27          c2 <-= (v - 2);

```


matchstmt:




Fig. 15. Syntax diagrams for match statements

```

28         r <-= <-c1 + <-c2;
29     }
30 };
31 }

```

Another implementation of a Fibonacci sequence generator namegen:

```

1
2 parfib : (int) : (int) =
3 {
4     c1, c2 : FibonacciType;
5
6     v := <-parfib;
7     matchseq
8     {
9         v == 0 => parfib <-= 0;
10        v == 1 => parfib <-= 1;
11        true  =>
12        {
13            match
14            {
15                true  => c1 = name2chan FibonacciType "parfib";
16                      c1 <-= (v - 1);
17
18                true  => c2 = name2chan FibonacciType "parfib";
19                      c2 <-= (v - 2);
20            }
21            parfib <-= <-c1 + <-c2;
22        }
23    }
24 }

```

In the above example, the inner `match` statement is used to initiate two Fibonacci number computations without restriction on the ordering (i.e., they could potentially execute in parallel).

3.11 The iter construct

The `iter` statement is a repetitive collection of guarded statement blocks. It executes repeatedly while any of its guards, which are Boolean expressions, evaluate to true. The `iter` statement is analogous to *guarded iteration* in Dijkstra's guarded commands [12]. All statements whose guards evaluate to true execute. The syntax of `iter` statements is illustrated by the syntax diagrams in Figure 16.

iterstmt:




Fig. 16. Syntax diagrams for iter statements

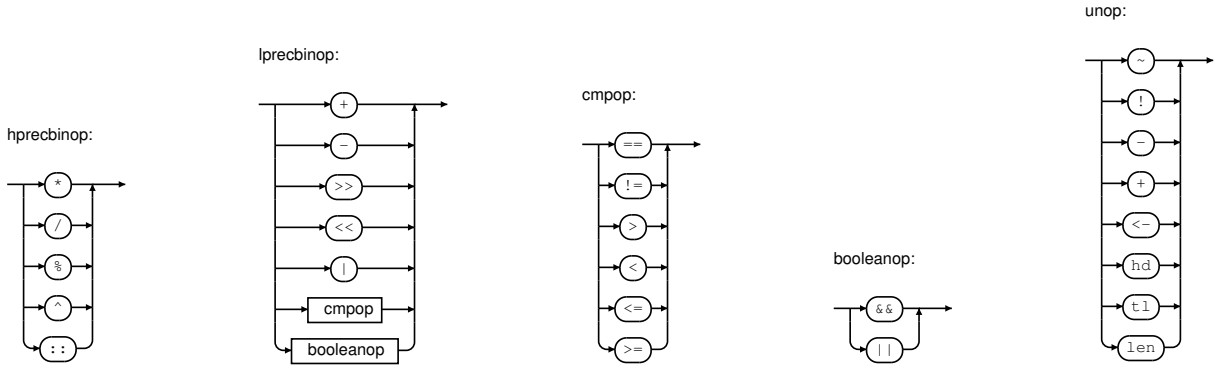


Fig. 17. Syntax diagrams for operators, and their classifications.

4. OPERATOR DESCRIPTIONS

The semantics of the language-level operators are presented in the following sections. The language includes operators for operation on *values*, operators on *names* and operators on *channels*. The language-level operators and their semantics are listed in Table 1 and Figure 17.

Table 1. Noisy language operators.

Operator	Description	Operator	Description
.	ADT member access	> =	Greater than or equal to
//	Array or character string subscript	= =	Equals
!	Logical not	! =	Not equals
~	Bitwise not	&	Bitwise AND, set intersection
++	Increment	^	Bitwise XOR
--	Decrement		Bitwise OR
hd	List head value	::	List append
tl	List tail value	&&	Logical AND
len	List, array, or string length; set size		Logical OR
name2chan	Bind name in name space to channel	=	Assignment
chan2name	Make channel visible as name in name space	: =	Declaration and assignment
var2name	Make variable visible as a name in name space	+ =	Addition/concatenation/union assignment
type	Type cast	- =	Subtraction/difference and assignment
* / %	Multiplication, division, modulo	* =	Multiplication and assignment
+	Unary plus, addition, set union, string concatenation	/ =	Division and assignment
-	Unary minus, subtraction, set difference	% =	Modulo and assignment
<<	Logical left shift	& =	Bitwise AND and assignment
>>	Logical right shift	=	Bitwise OR and assignment
<	Less than	^ =	Bitwise XOR and assignment
>	Greater than	<< =	Logical left shift and assignment
< =	Less than or equal to	>> =	Logical right shift and assignment
		<-	Assignment to/from channel

4.1 Operators on names

The operators on names are `name2chan`, `chan2name`, and `var2name`. The `name2chan` operator applied to a name (string constant) yields a channel to an entry in the name space having the same type as the *l-value* of the statement, i.e., it matches on both the name as well as type of an entry in the runtime name space. If the name corresponds to a `namegen`, a new instance of the `namegen` is created, as described in Section 6. If multiple matches exist, the match is non-deterministic. A `name2chan` expression completes as soon as it either succeeds or times out. If no match exists (timed out), the result of the `name2chan` expression is `nil`. Example:

```

1
2 #      Bind channel 'c' to the name "Fibonacci" having
3 #      type FibonacciType, timeout after 4 seconds
4 c      := name2chan FibonacciType "Fibonacci" 4E6;

```

New entries in the runtime name space can be synthesized as interfaces to channels or variables within a namegen, via the `chan2name` and `var2name` operators. If there is an extant entry in the name space of the parent proctype with the same name *and* type, it is replaced with the new entry, but only within the scope of the initiating namegen and namegens which it has instantiated via `name2chan`. When a variable which has been made visible in the name space via `var2name` is written to, the write not only updates the memory cell on the local device, but also generates a write to any channel bound to the name.

4.2 Operators on channels

There are two operators on channels, a channel send (`channel <:= expr`) and a channel receive (`variable or nil <- channel`). Channel communication is unbuffered, rendezvous: a send (receive) completes when a matching receive (send) is performed at the other end of the channel.

All interaction between namegens is over a channel. Since namegens might be executing locally or remotely, the passage of references down a channel is not permitted. E.g., a channel cannot be passed down a channel. Instead, the reference must be converted to a name via `chan2name`. The value of `chan2name` and `var2name` expressions is the name (type `string`) that was posted for the chan, or `nil` if the posting failed. The `chan2name` and `var2name` operators can either be used to post an explicitly chosen name, e.g.

```

1      chan2name mychan "name";

```

or might be used to generate an implicit name, in which case the name posted is chosen by the runtime and returned as the values of the `chan2name` or `var2name` expression.

5. TYPES

The syntax of type declarations and definitions were previously provided in Section 3. This section provides further detail on the language base types, as well as the construction of collection types.

The language includes a small set of basic types—`bool`, `nybble`, `byte`, `int`, `fixed` (fixed-point approximate real numbers), `real` (floating-point approximate real numbers) and `string`. The arithmetic types are `bool`, `nybble`, `byte`, `int`, `fixed` and `real`. The basic language data types may be employed in aggregate collections. The collection data types are arrays, tuples, ADTs (aggregate data types), lists and sets. Channels may be defined to carry any of the basic or aggregate data types, but these cannot include channels. Type definitions for the language base types may have an optional *error tolerance constraint*.

In addition to the basic types that can be used in type expressions, an identifier can also be of types `proctype`, `namegen` or act as an alias for a type expression. In the cases of `proctype` and `namegen`, the type signatures are functions of the body of the `proctype` definition, and that of the `namegen` channel interface, respectively. ADTs and type aliases introduce new *type names*.

5.1 Reference versus value types, channels and nil

All the primitive and aggregate types of variables are value types. Channels are reference types (they have *implicit state*), and may not be passed down channels. A primitive or aggregate variable cannot be transmitted down a channel if it contains a channel within it. Since the equivalent of function calls are communications on channels, and as previously stated, these communications can only carry values, it could be said that the semantics for interaction between program components are *call-by-value*. Treating all types as value types should not be thought of as an inefficiency burden as an implementation can use a copy-on-write strategy to make computation within a namegen efficient.

5.2 Specifying tolerance constraints

For some variables in programs, it might be known that a certain degree of deviation from their correct values is acceptable. In the case of channels, a certain degree of erasures might be tolerable. Tolerance constraints on variables

and channels, is key to being able to use the mathematical analysis of error magnitudes in programs [27] to implement forward error correction on the values exchanged between namegens. Conversely, it enables deliberate changes in functionality that improve performance but might induce errors.

For example, values to be transmitted on a communication channel might be occasionally purposefully dropped to reduce power consumption or improve *overall* network performance, while staying within prescribed constraints. As another example, in an application partitioned over devices that communicate over a wireless medium, the transmit power consumption might be purposefully reduced at the cost of a higher (tolerable) bit error rate. The low-level communication encodings are thus tuned to *application error tolerance properties*. We could think of the error tolerance constraints exposing semantic information that enables application-specific source and/or channel coding.

The most flexible form in which constraints specifying *tolerable error* could be provided, would be as a *tail distribution* on error magnitude. For example, it might be desirable that the probability that the magnitude of error in a variable is greater than x should vary with x as $\frac{1}{x}$. The notational complexity of representing such error tolerance constraints would be significant: since it is logical for the error tolerance constraint to be placed in the type annotation, the type would then need to contain an expression in a variable (x in the above example). This approach is therefore avoided. Experience with the language and error tolerance constructs may necessitate revisiting this restriction.

Instead of such general error tolerance constraints, the language design includes a restricted form of the above in which x is constant. An error tolerance constraint is thus defined in terms of two variables, m and A , specifying that the probability the magnitude of error in a variable is greater than m , should be less than A , i.e., $\Pr\{M > m\} \leq A$. A type expression involving the basic types can be qualified with one or more error, erasure or latency tolerance constraints. For example, the following declares a variable `pixel`, of type `(byte, byte, byte, byte)`, with the error constraint that the probability that the magnitude of error (due to erasures or errors) in a member of the 4-tuple exceeds 4 should be 0.01:

```

1
2 t      : type int epsilon(4, 0.01);
3 pixel  : (t, t, t, t);

```

Latency tolerances are inherently at odds with error and erasure tolerances. With errors and erasures, one can introduce redundancy in the form of encoding to mitigate the effects of the undesirable phenomena. In the case of latency tolerances on the other hand, one must reduce redundancy to improve (decrease) latency (transmission time). On the other hand, the violation of latency tolerances can be easily checked in an implementation (e.g., via timers), whilst the violation of error magnitude tolerances is more involved: violation of the tolerance inherently means the decoding process wrongly decoded an erroneous codeword because it had more errors than that which could be corrected based on the minimum Hamming distance in the chosen code³.

The goal of error and erasure tolerances in the current implementation is to address the problem of errors and erasures in the communication links between a partitioned application (between its constituent namegens). Tolerance constraints specified for channels apply to the values communicated on the channel. Tolerance constraints specified for variables are not used to encode the variables, but rather propagated upward through the dataflow graph to channels, reads from which reach the variables with tolerance constraints. Although a bit more difficult to reason about, the encodings also benefit values in programs, not only values communicated over a network (over channels): all variables can be thought of as channels to memory. This semantically exposes the fact that a read or a write from a “variable” might either mean reading from a local on-chip memory, or from a memory across the network, and the error tolerance associated with that “variable” might therefore either be used for bus encoding for error detection or forward error correction to a local memory, or for encoding a communication that happens over the network. A future implementation of the compiler could use the tolerance information on variables to encode the variables themselves, in order to mitigate the effect of soft-errors in memories, in an application-specific manner.

³The minimum distance of a code, $d(C) \geq \delta \iff C$ is $\delta - 1$ error correcting. $d(C) \geq 2\epsilon + 1 \iff C$ is ϵ error correcting [6].

5.3 Types bool, nybble, byte, int, real and fixed

The primitive types `bool`, `nybble` and `byte` are unsigned quantities with sizes 1, 4 and 8 bits respectively. The type `int` is a signed 32-bit value in two's complement format. The fixed point approximate real numbers, `fixed`, are signed 16-bit binary fixed point values in which 0 up to 15 bits can be used to represent the fractional component, and one bit is used to represent the sign. The floating point approximate real numbers, `real` are in 64-bit IEEE-754 double precision floating point format.

5.4 Type string

Strings are sequences of Unicode characters. An element in a string (a character), obtained using the notation `string_identifier[index]` is a 16-bit value. Strings can be concatenated using the `+` operator. The `len` operator, applied to a string yields the number of characters in the string. A substring or *slice* of a string is obtained using the notation:

`string_identifier[start index (optional) ':' end index (optional)]`

and yields a string. When the start (end) index is absent, it is implicitly the first (last) element of the string. The empty string `""` is represented by the value `nil`. Examples:

```

1
2 #       Define a new string 'name'
3 phrase  := "Mountains made of steam";
4
5 #       The variable c has type 'int', and holds the character 'M'
6 c       := phrase[0];
7
8 #       This sets the variable 'q' to "made of steam"
9 q       := phrase[len "Mountains ": ];

```

5.5 The array collection type

Arrays hold sequences of items of a single type. Multi-dimensional arrays are stored in memory in row-major order. The storage for arrays is conceptually allocated at the point of definition, thus there is no distinction between array declaration and array definition as there does in some languages. An element in an array is obtained using the notation `array_identifier[index]`. The `len` operator applied to an array yields the length of the array. A *slice* of an array is obtained using the notation:

`array_identifier[start index (optional) ':' end index (optional)]`

and yields an array of the same type. The type of an array includes its length, and thus all declarations of arrays include a specification of their length. However in an array definition from initializers, the dimensions may be omitted as they are implicit in the initializer. Example:

```

1
2 #       Declare 1-dimensional array 'a' of nybbles with 32 elements
3 a       : array [32] of nybble;
4
5 #       Declare a 3-dimensional array 'volume' of bits, and set one point
6 volume  : array [128][128][128] of bool;
7 volume[0][0][1] = 1;
8
9 #       Define a 3-dimensional array 'volume' of bits
10 volume2 := array [128] of { * => array [128] of { * => array [128] of bool } };
11
12 #       Definition of an array, type implicit from initializers,
13 firstArray := array of {"oil", "and", "water", "don't", "mix"};
14
15 #       Definition of an array, type implicit from initializers, with size 32
16 secondArray := array of {"oil", "and", "water", "don't", "mix", 31 => "."};
17
18 #       Definition of an array, type implicit from initializers, size 32;

```

```

19 #       the last 27 entries contain "."
20 thirdArray := array [32] of {"oil", "and", "water", "don't", "mix", * => "."};

```

5.6 Tuple collection type

A tuple is an unnamed collection of items of possibly different type. A tuple type might be considered as a cartesian product of type expressions. The elements in a tuple cannot be accessed individually; assignment to a tuple must be from a tuple expression, and assignment from a tuple type must be into a tuple of variables. Thus any assignment to a tuple, sets all fields. Example:

```

1
2 #       Declare color to be a tuple type
3 color   : (byte, byte, byte, byte);
4
5 color = (0, 6, 32, 8rFF);

```

5.7 Aggregate Data Type (ADT) collection type

ADTs are similar to tuples, with the only difference being that the members of an ADT may be named. An ADT declaration introduces a new type name, and subsequently, variables may be declared with that type name. Assignments from tuples with the same order and type of variables are permitted to ADT instance variables. Example:

```

1
2 #       Define an ADT type Color
3 Color : adt
4 {
5     r : byte;
6     g : byte;
7     b : byte;
8     a : byte;
9 };
10
11 #       Declare a variable with type Color
12 c : Color;
13
14 #       Assign a tuple to the ADT instance
15 c = (0, 6, 32, 8rFF);

```

5.8 The set collection type

The **set** collection data type is used to represent unordered collections of data items. A set can be cast to a list of the same type, yielding a non-deterministic ordered list. A list can likewise be cast as a set. The operations defined on set variables are *set union* ($A \cup B$ is denoted by $A + B$), *set intersection* ($A \cap B$ is denoted by $A \& B$), *set difference or relative complement* ($A \cap \overline{B}$ is denoted by $A - B$) and *set cardinality* ($|S|$ is denoted by $\text{len } S$). The type of a set includes its cardinality, and thus all declarations of sets include a specification of their cardinality. However in a set definition from initializers, the cardinality may be omitted as it is implicit in the initializer. Multiplicity of elements is ignored (i.e., sets are not *multisets*). An empty set has the value `nil`. Examples:

```

1
2 #       Declare 's' as a set with cardinality 32, of integers
3 s       : set [32] of int;
4
5 #       Add a few elements to the set
6 s += set of {2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41};
7
8 #       Check for set membership
9 u7p := s & set of {7};

```

5.9 Recursive list collection type

Recursive lists are lists of elements of a single type. The operations on `list` instances are `hd` (*head*), which yields single datum, the first element in the list, `tl` (*tail*), which yields a list representing all but the first element of the list, and `::` (*cons*), which is used to append an element to the head of the list. Examples:

```

1
2 #       Define 'veggies' as a list of strings
3 veggies := list of {"carrot", "celery", "radish"};
4
5 #       'celery' will be a variable of type string initialized to the string "celery"
6 celery := hd tl veggies;
7
8 #       The cons (::) operator adds an item to a list
9 together := "rabbits" :: hd veggies :: nil

```

5.10 Dynamic data structures

There are no pointers *per se* in Noisy: this restriction ensures it is always possible to partition programs by using the runtime name space as an interface between partitions.

The only new instances that can be created at runtime are new instances of namegens. Namegens which embody data structures can be used to achieve dynamic data structures. A point of elegance in this approach is that, in much the same way that namegens may be instantiated on any hardware device on the network, the data structures embodied in namegens may reside on arbitrary devices in the network. As an example, consider a data structure used to implement nodes in the abstract syntax tree of a parser, implemented as a namegen, where `Node` is a previously defined ADT:

```

1
2 node : (Node):(Node) =
3 {
4     n : Node;
5
6     iter
7     {
8         match
9         {
10             <-node => node <-= n;
11
12             node <- => n = <- node;
13         }
14         true => ;
15     }
16 }
17
18 creategraph : ():() =
19 {
20     n : Node;
21
22     n.left = name2chan "node";
23     n.right = name2chan "node";
24 }

```

Each time it is desired to dynamically create a new instance of the `Node` structure, a `name2chan` on the `node` namegen will create a new instance of the data structure, and the executor of the statement will obtain a channel to this new instance. In a sense, this approach bundles dynamically created data structures with their own access methods (through channel reads and writes). As a result, such data structures have a transactional interface.

5.11 The vector type

Vectors in Noisy are not collection types, but rather are vectors in the mathematical or physics sense: geometric objects in an n -dimensional space and denote a direction.

```

1
2 #       Declare a Cartesian vector
3 a       : vector [3] of real;
```

6. THE RUNTIME SYSTEM

6.1 Runtime system abstraction

One way to look at the runtime system is by analogy to the memory system in a stored program computer. In traditional models of a stored program computer (e.g., the von Neumann / Princeton model), programs access *memory cells* by performing operations on *addresses*. These addresses form an *address space*. Two operations can be performed on a memory cell, given an address that represents it: **memory read** and **memory write**. When programs in a language such as C allocate memory dynamically, they are essentially communicating with the operating system to obtain a new address. We could call this operation on the address space **new address**, although it doesn't really *introduce* a new address, but just provides access to the address, to a program. We can thus say that the alphabet of operations on memories is {**memory read**, **memory write**, **new address**}. In a computer system, these operations are manifest as operations in hardware: a **memread** operation is embodied in instructions in the machine's instruction set architecture which perform a read from memory, e.g., an **ld** (load) instruction. Furthermore, the execution of such an instruction leads to a communication between the processor and the memory, yielding the result of the load.

By analogy, the runtime system in Noisy is made up of *names*, which are accessed through *channels*. The alphabet of operations pertinent to names is {**nameread**, **namewrite**, **name2chan**, **chan2name**, **var2name**}. The **name2chan** operation obtains a new channel through which programs can write to an entry in the name space. Conversely, **chan2name** (**var2name**) makes a channel (variable) in a program visible as an entry in the name space. Entries in the name space have the type structure of the variables or channels with which they are associated. As an example, a **nameread** operation is logically associated with a channel read expression in a program, just as a **memread** in the discussion above is logically associated with reading the value of a variable. Just as a **memread** leads to a communication on the address and data buses to memory, a **nameread** causes a device on which a namegen is executing to communicate with other computing devices which are part of the name space. It is at the level of these messages that error, loss (erasure) and latency tolerance constraints for channels are used to drive encoding. It is also from the framework implementing these low level messages that the tolerance constraint violations on channels are obtained. The nature and protocol of these communications is the subject of Section 6.5.4.

6.2 Target hardware model

The underlying hardware *abstraction* is assumed to be a collection of computing devices. More concretely, it is assumed that each namegen is mapped onto a *virtual processing device*, and the virtual processing devices are interconnected via a single *virtual network*. The virtual processing devices might have a one-to-one or many-to-one correspondence with actual processing devices (e.g., microcontrollers, processor cores, microprocessors, workstations). Similarly, the virtual network might be composed of a multi-hop wired or wireless network, on-chip bus, on-chip network, etc. The virtual processing devices will subsequently be referred to as *processing devices*, and the virtual network as *the network*, without discussing further details of either the nature of processing devices, or issues such as the underlying implementation of device-to-device communication, routing within the network, or optimization of placement/mapping of namegens to actual physical processing devices. The combination of virtual processing devices and the virtual network will be referred to collectively as the *runtime system*. The term *loading a program* will subsequently be used to refer to the act of placing the compiled Noisy program on a virtual device. It should be stressed that the above statements describe an *abstraction*; their purpose is to enable a description of the system without getting mired in implementation-specific details.

6.3 Structure of the runtime name space

All namegens exposed in an application’s proctype definition appear as entries in the runtime name space with names of the form *proctype.namegen*. Entries in the runtime name space with distinct types associated with them are distinct. Duplicate entries (with identical names and types) may exist; this introduces nondeterminism.

A namegen is only *executing*, after its interface name has been bound to a channel. Each binding of a name to a channel creates a logically new instance of the namegen. Such a namegen instance so created shares the same name space as the namegen that bound its interface to a channel, and any changes to the runtime name space (e.g., creation of new entries via *chan2name* or *var2name*) are mutually visible. Besides entries corresponding to namegens, namegens may dynamically create entries in the runtime name space tied to channels. Such entries have the form *proctype.namegen.name*.

6.4 Enforcing access restrictions on names

The channels used by namegens to access names, can be used to implement the concept of *capabilities* from capability-based systems [21]. They provide an access path to software or hardware that is represented by a name in the name space. Access control can trivially be implemented by requiring the first value written down a channel to be an access right (e.g., a password)⁴.




Fig. 18. Structures underlying the implementation of the runtime system.

6.5 Runtime system implementation

The *runtime system* is the collection of data structures, instantiated at runtime, that support the underlying operations being performed by executing programs. At the heart of the runtime system implementation is a set of three tables: the *name table*, *activation record table* and the *chan table*, illustrated in Figure 18. The implementation of these tables is described in Section 6.5.1.

6.5.1 The name table. The name table contains an entry for each namegen installed on a device, along with an entry representing the type structure of the namegen. The *Name* entries are strings representing the namegen, qualified by the proctype of which they are part (entries in the runtime name space are of the form *proctype.namegen[.name]*). At runtime, new entries are added to the name table whenever an executing namegen performs a *chan2name* or *var2name* operation. New entries in the name table are created whenever new code is installed on a device. By convention, a namegen with the identifier *init* immediately begins executing once loaded. Loading a namegen implementation with the same proctype as an extant one, into the runtime system (e.g., on a different device) should be thought of as being the same as overwriting code memory of a running application in a traditional program; undefined behavior will result.

When a namegen performs a *name2chan* operation, the local name table is first consulted. If no matching name and type is found locally, the name tables of all devices in the network are consulted⁵. If such an operation is successful,

⁴Furthermore, a particular application might even define the sequence of accesses that must precede “real” interaction, to engage the accessor in, say, a Diffie-Hellman key generation, generating a shared secret that can be used to encrypt subsequent communications.

⁵Logically, the query is a broadcast, but an implementation may perform any number of optimizations to make this lookup more efficient.

i.e., the name and type match an entry in the local name table, a new instance of the namegen (based on the PC entry) begins executing, with its own private stack. Such an instance is termed a namegen *activation*. The state for currently instantiated namegens is maintained in the *activation record table*.

6.5.2 The activation record table. The activation record table maintains the state corresponding to each namegen instantiation (as created by a `name2chan`, or an `init namegen` (Section 6.5.1)). The `ID` field uniquely identifies an instantiated namegen, and is used to identify namegens for all other operations. For example, all channels are associated with a particular namegen instance, and the instance's `ID` is used to track this correspondence. An instantiation of a namegen may create new entries in the runtime name space; these are only visible to the namegen that caused their instantiation (and to themselves).

6.5.3 The chan table. The chan table contains entries for all channels that are within the current scope of execution of all namegen activations. A namegen that performs a send or receive operation on a channel sleeps on a *rendezvous structure* in the chan table. When the channel communication operation completes (e.g., message successfully transmitted over network and an acknowledgment received), the sleeping namegen is awoken. The `xform` field contains a matrix (logically a part of the channel's type) representing the linear transformations that must be used to encode and decode the data exchanged between devices, representing the communication on a language-level channel. The messages on the network which are generated as a result of operations on channels are described in the next Section.

Table 2. Name communication protocol.

Message	Description	Associated Language Construct	Parameters
Tname2chan	Bind name to channel; if name is a namegen, create instance	name2chan	name, type
Rname2chan	Response: chan ID or nil		
Tnameread	Channel receive	Channel receive expression (<code><-c</code>)	chan ID
Rnameread	Response: type structured data		
Tnamewrite	Channel send	Channel send expression (<code>c <-=</code>)	data, chan ID
Rnamewrite	Acknowledgment		

6.5.4 Name communication protocol. A small alphabet of messages may be exchanged between devices as a result of language-level constructs related to channels. The list of messages in this alphabet is provided in Table 2. The following detail the effect of the receipt of messages in Table 2, on the runtime table structures, and on execution at the recipient node. The only language level channel operators that do not lead to the generation of messages are `chan2name` and `var2name`. They lead only to the creation of entries in the local name table.

Tname2chan / Rname2chan. Execution of a `name2chan` expression in a namegen will initiate the generation of a `Tname2chan` message on the network. A device which contains a matching entry (on both name and type) in its name table responds with a `Rname2chan` message. It is possible that no such device might exist, in which case the language-level expression will evaluate to `nil` after the language-level-specified timeout.

If the supplied name, in the name space, is a namegen, a new activation of the remote namegen is created, and corresponding entries for the send / receive interface channel tuple in the remote device's chan table are created. The ID of this allocated entry is returned to the initiating node in a `Rname2chan` message. A new entry is created in a local *chan table*, and this entry is used to store the received tag. Subsequent operations on the channel associated with this `name2chan` operation will occur with the specific instantiation of the remote namegen.

Tnamewrite / Rnamewrite. A channel send operation causes a `Tnamewrite` message to be generated on the network. The message target is determined by a lookup in the local chan table for (1) the destination network address and (2) the destination namegen ID, and (3) the destination namegen channel index, which identifies a channel in a specific instance of the remote namegen (with its private execution stack) that the values should be

delivered to. The `timer` entry of the local `chan` table is updated with a timestamp, which is used to determine timeouts when the channel appears in a latency violation expression.

Tnameread / Rnameread. A channel receive operation causes a **Tnameread** message to be generated on the network. A lookup is performed in the local `chan` table to determine the remote device and the ID, to be included in the **Tnameread** message, designating which channel in a specific instantiation of a remote `namegen` the message should be delivered to. The `timer` entry of the local `chan` table is updated with a timestamp, which is used to determine timeouts when the channel appears in a latency violation expression.

APPENDIX

A. A CALL-BY-VALUE LAMBDA CALCULUS WITH ERROR-TOLERANT TYPES

A central theme of the ideas presented in Noisy, is that, rather than adopting the traditional method of duplicating (or triplicating) computation in order to counteract the occurrence of failures, we introduce the concept of *error-tolerance constraints in programs*. To formalize this notion, we introduce a typed lambda calculus, λ_ε , in which the type ascriptions have error tolerance constraints. The role of λ_ε in this work is analogous to λ_{zap} in [34]. In contrast to [34], our goal is to enable programs to specify the amount of tolerable *numeric error magnitude*.

Type Inference Rules for λ_ε

$\frac{\text{T-INT}}{\Gamma \vdash n, \varepsilon : \text{int}, \varepsilon}$
$\frac{\text{T-TRUE}}{\Gamma \vdash \text{true}, \varepsilon : \text{bool}, \varepsilon}$
$\frac{\text{T-FALSE}}{\Gamma \vdash \text{false}, \varepsilon : \text{bool}, \varepsilon}$
$\frac{\text{T-CONSTRAINTPRESERVATIONUNDERERROR} \quad v : T_1, \varepsilon_1 \quad K_{\varepsilon, f_t} : v \rightarrow v'}{\Gamma \vdash^{(t) \sim f_t} v' : T, \varepsilon}$
$\frac{\text{T-ADD} \quad \Gamma \vdash v : T, \varepsilon_1 \quad \Gamma \vdash w : T, \varepsilon_2}{\Gamma \vdash v + w : T, q(\varepsilon_1, \varepsilon_2)}$
$\frac{\text{T-IF} \quad \Gamma \vdash^{(t) \sim f_t} \text{cond} : \text{bool}, \varepsilon \quad \Gamma \vdash^{(t) \sim f_t} a : T, \varepsilon_1 \quad \Gamma \vdash^{(t) \sim f_t} b : T, \varepsilon_2}{\Gamma \vdash^{(t) \sim f_t} \text{if } \text{cond} \text{ then } a \text{ else } b : T, q(\varepsilon_1, \varepsilon_2)}$
$\frac{\text{T-ABS} \quad \Gamma, x : T_1, \varepsilon_1 \vdash y : T_2, \varepsilon_2}{\Gamma \vdash \lambda x : T_1, \varepsilon_1. y : (T_1, \varepsilon_1 \rightarrow T_2, \varepsilon_2)}$
$\frac{\text{T-APP} \quad \Gamma \vdash g : (T_1, \varepsilon_1 \rightarrow T_2, \varepsilon_2) \quad \Gamma \vdash x : T_1, \varepsilon_1}{\Gamma \vdash g \ x : T_2, \varepsilon_2}$
$\frac{\text{T-LET} \quad v : T, \varepsilon_1 \quad w : T, \varepsilon_2}{\text{let } v = w \text{ in } e : T, r(\varepsilon_1, \varepsilon_2)}$

The type inference rules for λ_ε are shown above. The first three inference rules are straightforward. For example, the value `true` with error tolerance constraint ε has type `bool`, with type error constraint ε . The fourth type inference rule, T-CONSTRAINTPRESERVATIONUNDERERROR is the key component that captures our proposal of error-tolerance transformations (encodings). The concept it embodies is that, if v has type T , and error tolerance

constraint ε , and K_{ε, f_t} is a transformation that takes as parameters the error tolerance constraint ε and a bit upset probability distribution, f_t , and encodes v to give v' , then under the occurrence of a bit upset pattern $\langle t \rangle$, which follows distribution f_t , v' obeys the type and error constraint ascriptions of v .

B. FORMAL DEFINITION OF TYPE SYSTEM

The type inference rules specify the type resulting from the combination of expressions using the language level operators.

C. IMPLEMENTATION

Figure 22 shows a screenshot of the web interface to a prototype implementation of the compiler. This interface is also available online ([here](#)).

The Noisy language grammar is listed in Appendix D. The language design is influenced most directly by the C (statement syntax) and Limbo (e.g., `adt` and `list` data type) programming languages, Hoare’s CSP (channel concept) and Dijkstra’s guarded commands (the `match` and `matchseq` statements), among others. Many of these influences were themselves influenced by earlier languages — Algol in the case of C, CSP, Alef and Newsqueak in the case of Limbo, etc. The overall *programming model* that Noisy provides bears some similarities to the Actors model [1]. The role of the runtime system and its associated language constructs (`name2chan`, `chan2name`, `var2name`) pervades the language structure. The `name2chan` construct enables both the facilities akin to function calls / process creation (function calls and `spawn` in Limbo). The `chan2name` and `var2name` language constructs were inspired by experience with the `sys->file2chan()` system library routine in the Inferno operating system [26]. The language level error, erasure and latency tolerance constraints as well as the constraint-violation driven control flow were borne out of our investigation of language-level transformations for error tolerance, and the programming language constructs which support them [29].

The process of designing the grammar was driven by the desire to attain a language structure that was relevant to our goals of ease of partitioning and error-correctness tradeoffs, as well as a desire to simplify the implementation. The grammar is LL-1 to ensure that a recursive descent parser can be built for it, using only a single token of lookahead.

The compiler implementation was heavily influenced by Niklaus Wirth’s Oberon compilers. Parsing is implemented with a recursive descent parser. Unlike Wirth’s compilers which generate code in a single pass, the parser is used to build a tree-based intermediate representation, an *abstract syntax tree* (AST). The AST is an abbreviated form of the parse tree, with nodes whose presence is implicit in a given subtree removed. Our use of the AST is essentially as an intermediate representation⁶. This in-memory intermediate representation is then traversed by one of the subsequent stages. There are currently two second-stage passes that use this intermediate representation — graph output, and C code generation. The graph output pass walks the AST and symbol table data structures, and outputs it in *dot* format, which is passed to the Dot [13] tool to render the graphs in one of many formats, including Postscript. Figures C and 24 illustrate a small example program, its AST and the associated symbol table graphs generated by the compiler and rendered with Dot.

C.1 Building the AST

The lexical analyzer converts the input source into a sequence of *tokens*, corresponding to the reserved words, operators and separators in the language, and identifiers. The parser consumes these tokens from the lexer in the process of recognizing syntactic forms. The parser is structured as a set of routines, one per grammar production. Starting with the start symbol of a program, it recursively processes each valid production at that point in the parse. The decision of which production, and hence which routine to invoke next, is driven by the grammar $FIRST(X)$ and $FOLLOW(A)$ sets. During the language design, left recursion was eliminated from the grammar using Algorithm 4.1 in [2], and the grammar was left-factored using Algorithm 4.2 in [2]. The elimination of left-recursion is important as the $FIRST(X)$ sets of two alternatives in a grammar production may overlap in the presence of left-recursion. The top-level structure of the compiler is illustrate in Figure 21.

⁶Pascal P-code was essentially a linearized form of the AST.

The AST is a binary tree; when nodes in the tree have greater than two children, the children are hung off a subtree of nodes used for chaining. Among other things, nodes contain references to entries in the symbol table, and to canonical trees representing their types, if relevant.

C.2 Modifying the AST for easier code generation

There exist a few productions in the grammar which have as their result non-terminals. While this unclutters the grammar, it also leads in practice to AST subtrees that are unnecessarily “tall”. For example, the grammar production for a constant declaration is:

```
1 condecl ::= "const" (intconst | realconst | boolconst) .
```

The AST subtree built for a recognized `condecl` production will have as its root a node with type `condecl`, which will have as its only child (the token `const` is redundant in the AST) a node with type being one of `intconst`, `realconst` or `boolconst`. The subtree of height two can however be substituted directly with one of these children. For lack of a better term, we may call this “tree height shortening”. Tree shortening does not affect the functionality of the AST, and it greatly simplifies the task of routines that subsequently traverse the tree to generate code or perform other actions.

C.3 The symbol table

The symbol table stores context sensitive information gleaned from the input source during the process of parsing. Its structure is simple: it is a generalized tree with each subtree corresponding to a *scope*. Linked off each node in this tree is a list of identifiers, corresponding to the identifiers defined within that scope.

C.4 Types

Types are represented in the symbol table as trees built out of the primitive types and collection type constructors. These *type trees* correspond to the parse tree of a type expression. For progtypes, the ordering of definition of namegens does not affect the type signature, so there are multiple type trees for a given signature. In all other cases, two types are equivalent if their type trees are identical. Type equivalence is thus *structural* as opposed to *name equivalence*. To check types, the compiler performs a post-order walk of a type tree and generates a signature based on the nodes visited. Such a signature uniquely identifies a type. The graph generators include in the generated graphs textual renderings of these signatures on each node representing an identifier.

D. Noisy LANGUAGE GRAMMAR

```

1 character      ::= Unicode-0000h-to-Unicode-FFFFh .
2 rsvopseptoken ::= "~" | "!" | "%" | "^" | "&" | "*" | "(" | ")" | "," | "-" | "+" | "="
3               | "/" | ">" | "<" | ";" | ":" | "'" | "\"" | "{" | "}" | "[" | "]" | "|" |
4               | "<-" | "." | "<=" | ">=" | "^=" | "|=" | "&=" | "%=" | "/=" | "*=" | "-="
5               | "+=" | "!=" | "!=" | ">>" | ">>=" | "<<" | "<<=" | "<-" | "&&" | "||"
6               | ":" | ">=" | "==" | "++" | "--" | "<=" .
7 zeronine      = "0-9" .
8 onenine       = "1-9" .
9 radix         = onenine {zeronine} "r" .
10 charconst    = "'" character "'" .
11 intconst     ::= [radix] ("0" | onenine {zeronine}) | charconst .
12 boolconst    ::= "true" | "false" .
13 drealconst    = ("0" | onenine {zeronine}) "." {zeronine} .
14 erealconst    = (drealconst | intconst) ("e" | "E") intconst .
15 realconst    ::= drealconst | erealconst .
16 strconst     ::= "\" {character} "\" .
17 idchar       = char-except-rsvopseptoken .
18 identifier   ::= (idchar-except-zeronine) {idchar} .
19
20 program      ::= progtypedekl {namegendefn} .
21 progtypedekl = identifier ":" "proctype" "{" progtypebody "}" .
22 progtypebody = {ptypenamedekl ";" } .
23 ptypenamedekl ::= identlist ":" (condecl | typedekl | namegendekl) .
24 condecl      = "const" (intconst | realconst | boolconst) .
25 typedekl     = ("type" typeexpr) | adttypedekl | vectortypedekl .
26 adttypedekl  ::= "adt" "{" identlist ":" typeexpr ";" {identlist ":" typeexpr ";" } "}" .
27 vectortypedekl ::= "vector" "[" intconst "]" of typeexpr .
28 namegendekl  ::= "namegen" tupletype ":" tupletype .
29 identornil   ::= (identifier {fieldselect}) | "nil" .
30 identornillist = identornil {"," identornil} .
31 identlist    = identifier {"," identifier} .
32 typeexpr     ::= (basictype [tolerance {"", tolerance}]) | anonaggrtype | typename .
33 typename     ::= identifier ["->" identifier] .
34 tolerance    = errormagtolerance | losstolerance | latencytolerance .
35 errormagtolerance = "epsilon" "(" realconst "," realconst ")" .
36 losstolerance  = "alpha" "(" realconst "," realconst ")" .
37 latencytolerance = "tau" "(" realconst "," realconst ")" .
38 basictype     ::= "bool" | "nybble" | "byte" | "string" | "int" | realtype .
39 realtype      = "real" | fixedtype .
40 fixedtype     = "fixed" intconst "." intconst .
41 anonaggrtype  ::= arraytype | listtype | tupletype | settype .
42 arraytype     = "array" "[" intconst "]" {"[" intconst "]" of typeexpr} .
43 listtype      = "list" "of" typeexpr .
44 tupletype     = "(" typeexpr {"," typeexpr} ")" .
45 settype       = "set" "[" intconst "]" "of" typeexpr .
46 initlist     = "{" expr {"," expr} "}" .
47 idxinitlist   = "{" element {"," element} "}" .
48 starinitlist  = "{" element {"," element} ["*", ">=" expr] "}" .
49 element       = expr [ ">=" expr ] .
50 namegendefn   ::= identifier [":" tupletype ":" tupletype] "=" scopedstmtlist .
51 scopedstmtlist ::= "{" stmtlist "}" .
52 stmtlist     = {stmt} .
53 stmt         ::= [ identornillist (":" (condecl | typedekl | typeexpr)) | (assignop expr)
54               | "(" identornillist ")" assignop expr | matchstmt | iterstmt
55               | scopedstmtlist ] ";" .
56 assignop     ::= "=" | "^=" | "|=" | "&=" | "%=" | "/=" | "*=" | "-=" | "+=" | ">="
57               | "<<=" | "<-" | "!=" .
58 matchstmt    ::= ("match" | "matchseq") "{" guardbody "}" .

```

```

59 iterstmt      ::= "iter" "{" guardbody "}" .
60 guardbody     = {expr ">=" stmtlist} .
61 expr          ::= (term {lprecbinop term}) | anonaggrcastexpr | chanevtexpr
62              | chan2nameexpr | var2nameexpr | name2chanexpr .
63 listcastexpr  = "list" "of" initlist .
64 setcastexpr   = "set" "of" initlist .
65 arrcastexpr   = "array" (("of" idxinitlist) | ("[" intconst "]" of starinitlist)) .
66 anonaggrcastexpr ::= listcastexpr | setcastexpr | arrcastexpr .
67 chanevtexpr    ::= ("erasures" | "errors" | "latency") "of" identifier cmpop expr .
68 chan2nameexpr  ::= "chan2name" factor [strconst] .
69 var2nameexpr   ::= "var2name" factor [strconst] .
70 name2chanexpr  ::= "name2chan" typeexpr expr realconst .
71 term           ::= [basictype] [unop] factor ["++" | "--"] {hprecbinop factor} .
72 factor         ::= (identifier {fieldselect}) | intconst | realconst | strconst | boolconst
73              | "(" expr ")" .
74 fieldselect    ::= ("." identifier) | ("[" expr [":" expr "]" ) .
75 hprecbinop     ::= "*" | "/" | "%" | "^" | "::" .
76 lprecbinop     ::= "+" | "-" | ">>" | "<<" | "|" | cmpop | booleanop .
77 cmpop          ::= "==" | "!=" | ">" | "<" | "<=" | ">=" .
78 booleanop      ::= "&&" | "||" .
79 unop           ::= "~" | "!" | "-" | "+" | "<-" | "hd" | "tl" | "len" .

```

REFERENCES

- [1] G. Agha. An overview of actor languages. In *Proceedings of the 1986 SIGPLAN workshop on Object-oriented programming*, pages 58–67, New York, NY, USA, 1986. ACM Press.
- [2] A. V. Aho, R. Sethi, and J. D. Ullman. *Compilers: Principles, Techniques, and Tools*. Addison-Wesley, Reading, Massachusetts, 1986.
- [3] E. Allen, D. Chase, V. Luchangco, J.-W. Maessen, and G. L. Steele, Jr. Object-oriented units of measurement. In *Proceedings of the 19th Annual ACM SIGPLAN Conference on Object-oriented Programming, Systems, Languages, and Applications*, OOPSLA '04, pages 384–403, New York, NY, USA, 2004. ACM.
- [4] T. Antoniu, P. A. Steckler, S. Krishnamurthi, E. Neuwirth, and M. Felleisen. Validating the unit correctness of spreadsheet programs. In *Proceedings of the 26th International Conference on Software Engineering*, ICSE '04, pages 439–448, Washington, DC, USA, 2004. IEEE Computer Society.
- [5] M. Babout, H. Sidhoum, and L. Frecon. Ampere: a programming language for physics. *European Journal of Physics*, 11(3):163, 1990.
- [6] J. Baylis. *Error Correcting Codes: A Mathematical Introduction*. Chapman & Hall/CRC, 1997.
- [7] G. Biggs and B. A. Macdonald. A pragmatic approach to dimensional analysis for mobile robotic programming. *Auton. Robots*, 25(4):405–419, Nov. 2008.
- [8] R. D. Blumofe, C. F. Joerg, B. C. Kuszmaul, C. E. Leiserson, K. H. Randall, and Y. Zhou. Cilk: An efficient multithreaded runtime system. In *Proceedings of the Fifth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPOPP '95, pages 207–216, New York, NY, USA, 1995. ACM.
- [9] J. Bornholt, T. Mytkowicz, and K. S. McKinley. Uncertain: A first-order type for uncertain data. In *Proceedings of the 19th International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS '14, pages 51–66, New York, NY, USA, 2014. ACM.
- [10] F. Chen, G. Roşu, and R. P. Venkatesan. Rule-based analysis of dimensional safety. In *Proceedings of the 14th International Conference on Rewriting Techniques and Applications*, RTA'03, pages 197–207, Berlin, Heidelberg, 2003. Springer-Verlag.
- [11] R. F. Cmelik and N. H. Gehani. Dimensional analysis with c++. *IEEE Softw.*, 5(3):21–27, May 1988.
- [12] E. W. Dijkstra. Guarded commands, nondeterminacy and formal derivation of programs. *Commun. ACM*, 18(8):453–457, 1975.
- [13] E. R. Gansner and S. C. North. An open graph visualization system and its applications to software engineering. *Software — Practice and Experience*, 30(11):1203–1233, 2000.
- [14] P. N. Hilfinger. An ada package for dimensional analysis. *ACM Trans. Program. Lang. Syst.*, 10(2):189–203, Apr. 1988.
- [15] M. Hills, F. Chen, and G. Roşu. A rewriting logic approach to static checking of units of measurement in c. *Electron. Notes Theor. Comput. Sci.*, 290:51–67, Dec. 2012.
- [16] C. A. R. Hoare. Communicating sequential processes. *Commun. ACM*, 21(8):666–677, 1978.
- [17] R. T. House. A proposal for an extended form of type checking of expressions. *Comput. J.*, 26(4):366–374, Nov. 1983.
- [18] A. Kennedy. Dimension types. In *Proceedings of the 5th European Symposium on Programming: Programming Languages and Systems*, ESOP '94, pages 348–362, London, UK, UK, 1994. Springer-Verlag.
- [19] A. Kennedy. Types for units-of-measure in f#: Invited talk. In *Proceedings of the 2008 ACM SIGPLAN Workshop on ML*, ML '08, pages 1–2, New York, NY, USA, 2008. ACM.
- [20] A. J. Kennedy. Relational parametricity and units of measure. In *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '97, pages 442–455, New York, NY, USA, 1997. ACM.
- [21] H. M. Levy. *Capability-Based Computer Systems*. Digital Press, Bedford, Massachusetts, 1984.
- [22] J. Lim, P. Stanley-Marbell, and M. Rinard. Newton Language.
- [23] Motorola. Serial Peripheral Interface (SPI).
- [24] NXP Semiconductors. UM10204, *I2C-bus specification and user manual*, April 2014.
- [25] M. Rittri. Dimension inference under polymorphic recursion. In *Proceedings of the Seventh International Conference on Functional Programming Languages and Computer Architecture*, FPCA '95, pages 147–159, New York, NY, USA, 1995. ACM.
- [26] Sean Dorward, Rob Pike, David Leo Presotto, Dennis Ritchie, Howard Trickey and P. Winterbottom. The inferno operating system. *Bell Labs Technical Journal*, 2(1):5–18, Winter 1997.
- [27] P. Stanley-Marbell. Derivation of Probability Distributions of Data Values and Error Magnitudes in Program Variables Under Single and Multi-Bit Errors and Erasures (draft).
- [28] P. Stanley-Marbell, V. Estellers, and M. Rinard. Crayon: Saving power through shape and color approximation on next-generation displays. In *Proceedings of the Eleventh European Conference on Computer Systems*, EuroSys '16, pages 11:1–11:17, New York, NY, USA, 2016. ACM.

- [29] P. Stanley-Marbell and D. Marculescu. A Programming Model and Language Implementation for Concurrent Failure-Prone Hardware. In *Proceedings of the Workshop on Programming Models for Ubiquitous Parallelism*, September 2006.
- [30] P. Stanley-Marbell and M. Rinard. Lax: Driver interfaces for approximate sensor device access. In *15th Workshop on Hot Topics in Operating Systems (HotOS XV)*, Kartause Ittingen, Switzerland, May 2015. USENIX Association.
- [31] P. Stanley-Marbell and M. Rinard. Reducing serial i/o power in error-tolerant applications by efficient lossy encoding. In *Proceedings of the 53rd Annual Design Automation Conference*, DAC '16, pages 62:1–62:6, New York, NY, USA, 2016. ACM.
- [32] Z. D. Umrigar. Fully static dimensional analysis with c++. *SIGPLAN Not.*, 29(9):135–139, Sept. 1994.
- [33] Unicode Consortium. *The Unicode Standard: Version 4.0*. Addison Wesley, Reading, Massachusetts, Aug. 2003.
- [34] D. Walker, L. Mackey, J. Ligatti, G. Reis, and D. August. Static typing for a faulty lambda calculus. In *ACM SIGPLAN International Conference on Functional Programming*, New York, NY, USA, September 2006. ACM Press.
- [35] N. Wirth. What can we do about the unnecessary diversity of notation for syntactic definitions? *Commun. ACM*, 20(11):822–823, 1977.
- [36] N. Wirth. *Grundlagen und Techniken des Compilerbaus*. Addison-Wesley, Bonn, 1 edition, 1996.

Type Inference Rules for Operators

$\frac{\text{T-TRUE}}{\text{true} : \text{bool}}$	
$\frac{\text{T-FALSE}}{\text{false} : \text{bool}}$	
$\frac{\text{T-VAR}}{x^\tau : \tau}$	
$\frac{\text{T-ADTMEMBER}}{M : (\tau_1 \times \tau_2 \times \dots \times \tau_n)} \quad M.k : \tau_k$	
$\frac{\text{T-ARRAYELEM}}{M : (\tau \times \tau \times \dots \times \tau)} \quad M[k] : \tau$	
$\frac{\text{T-LOGICALNOT}}{M : \text{bool}} \quad !M : \text{bool}$	
$\frac{\text{T-BITWISENOT}}{M : \tau \quad \tau : \text{bool nybble byte int}} \quad \sim M : \tau$	
$\frac{\text{T-INCREMENT}}{M : \tau \quad \tau : \text{bool nybble byte int real fixed}} \quad M++ : \tau$	
$\frac{\text{T-DECREMENT}}{M : \tau \quad \tau : \text{bool nybble byte int real fixed}} \quad M- : \tau$	
$\frac{\text{T-HEAD}}{M : \text{list of } \tau} \quad \text{hd } M : \tau$	
$\frac{\text{T-TAIL}}{M : \text{list of } \tau} \quad \text{tl } M : \text{list of } \tau$	
$\frac{\text{T-LENGTH}}{M : \tau \quad \tau : \text{array string list set}} \quad \text{len } M : \text{int}$	
$\frac{\text{T-NAME2CHAN}}{M : \text{string}} \quad \text{name2chan } M : \text{chan of } \tau$	
$\frac{\text{T-CHAN2NAME}}{M : \text{chan of } \tau} \quad \text{chan2name } M : \text{string}$	
	$\frac{\text{T-VAR2NAME}}{M : \text{chan of } \tau} \quad \text{var2name } M : \text{string}$
	$\frac{\text{T-MULTIPLICATION}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real}} \quad M_1 * M_2 : \tau$
	$\frac{\text{T-DIVISION}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real}} \quad M_1 / M_2 : \tau$
	$\frac{\text{T-MODULO}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real}} \quad M_1 \% M_2 : \tau$
$\frac{\text{T-ADD}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real set string}} \quad M_1 + M_2 : \tau$	
$\frac{\text{T-SUB}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real set}} \quad M_1 - M_2 : \tau$	
	$\frac{\text{T-LSHIFT}}{M_1 : \tau \quad \tau : \text{nybble byte int fixed real} \quad M_2 : \text{int}} \quad M_1 \ll M_2 : \tau$
	$\frac{\text{T-RSHIFT}}{M_1 : \tau \quad \tau : \text{nybble byte int fixed real} \quad M_2 : \text{int}} \quad M_1 \gg M_2 : \tau$
	$\frac{\text{T-LESSTHAN}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real string}} \quad M_1 < M_2 : \text{bool}$
	$\frac{\text{T-GREATERTHAN}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real string}} \quad M_1 > M_2 : \text{bool}$
	$\frac{\text{T-LESSTHANEQUALS}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real string}} \quad M_1 \leq M_2 : \text{bool}$
	$\frac{\text{T-GREATERTHANEQUALS}}{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool nybble byte int fixed real string}} \quad M_1 \geq M_2 : \text{bool}$

Fig. 19. Type inference rules for language operators (part 1).

Type Inference Rules for Operators (*continued*)

T-EQUALS $\frac{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool} \text{nybble} \text{byte} \text{int} \text{fixed} \text{real} \text{string}}{M_1 == M_2 : \text{bool}}$	
T-NOTEQUALS $\frac{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool} \text{nybble} \text{byte} \text{int} \text{fixed} \text{real} \text{string}}{M_1 != M_2 : \text{bool}}$	
T-BITWISEAND $\frac{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool} \text{nybble} \text{byte} \text{int}}{M_1 \& M_2 : \tau}$	
T-SETINTERSECTION $\frac{M_1 : \text{set of } \tau \quad M_2 : \text{set of } \tau}{M_1 \& M_2 : \text{set of } \tau}$	
T-BITWISEXOR $\frac{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool} \text{nybble} \text{byte} \text{int}}{M_1 \wedge M_2 : \tau}$	
T-BITWISEOR $\frac{M_1 : \tau \quad M_2 : \tau \quad \tau : \text{bool} \text{nybble} \text{byte} \text{int}}{M_1 M_2 : \tau}$	
T-LISTCONS $\frac{M_1 : \tau \quad M_2 : \text{list of } \tau}{M_1 :: M_2 : \text{list of } \tau}$	
T-LOGICALAND $\frac{M_1 : \text{bool} \quad M_2 : \text{bool}}{M_1 \&\& M_2 : \text{bool}}$	
T-LOGICALOR $\frac{M_1 : \text{bool} \quad M_2 : \text{bool}}{M_1 M_2 : \text{bool}}$	
T-ASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 = M_2 : \tau}$	
T-DECLASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 := M_2 : \tau}$	
T-ADDASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 += M_2 : \tau}$	
T-SUBASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 -= M_2 : \tau}$	
T-MULASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 *= M_2 : \tau}$	
T-DIVASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 /= M_2 : \tau}$	
T-MODASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 \% = M_2 : \tau}$	
T-ANDASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 \&= M_2 : \tau}$	
T-ORASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 = M_2 : \tau}$	
T-XORASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 ^= M_2 : \tau}$	
T-LSHIFTASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 \ll = M_2 : \tau}$	
T-RSHIFTASSIGNMENT $\frac{M_1 : \tau \quad M_2 : \tau}{M_1 \gg = M_2 : \tau}$	
T-CHANREADTOLVALUE $\frac{M_1 : \tau_1 \quad M_2 : \text{chan of } \tau_1 \quad M_3 : \text{chan of } \tau_2}{M_3 = M_1 = <- M_2 : \text{chan of } \tau_2}$	
T-CHANREADTEST $\frac{M_2 : \text{chan of } \tau}{<- M_2 : \text{bool}}$	
T-CHANWRITETEST $\frac{M_2 : \text{chan of } \tau}{M_2 <- : \text{bool}}$	
T-CHANWRITE $\frac{M_1 : \tau_1 \quad M_2 : \text{chan of } \tau_1 \quad M_3 : \text{chan of } \tau_2}{M_3 = M_2 <- = M_1 : \text{chan of } \tau_2}$	

Fig. 20. Type inference rules for language operators (part 2).




Fig. 21. Structure of the compiler.

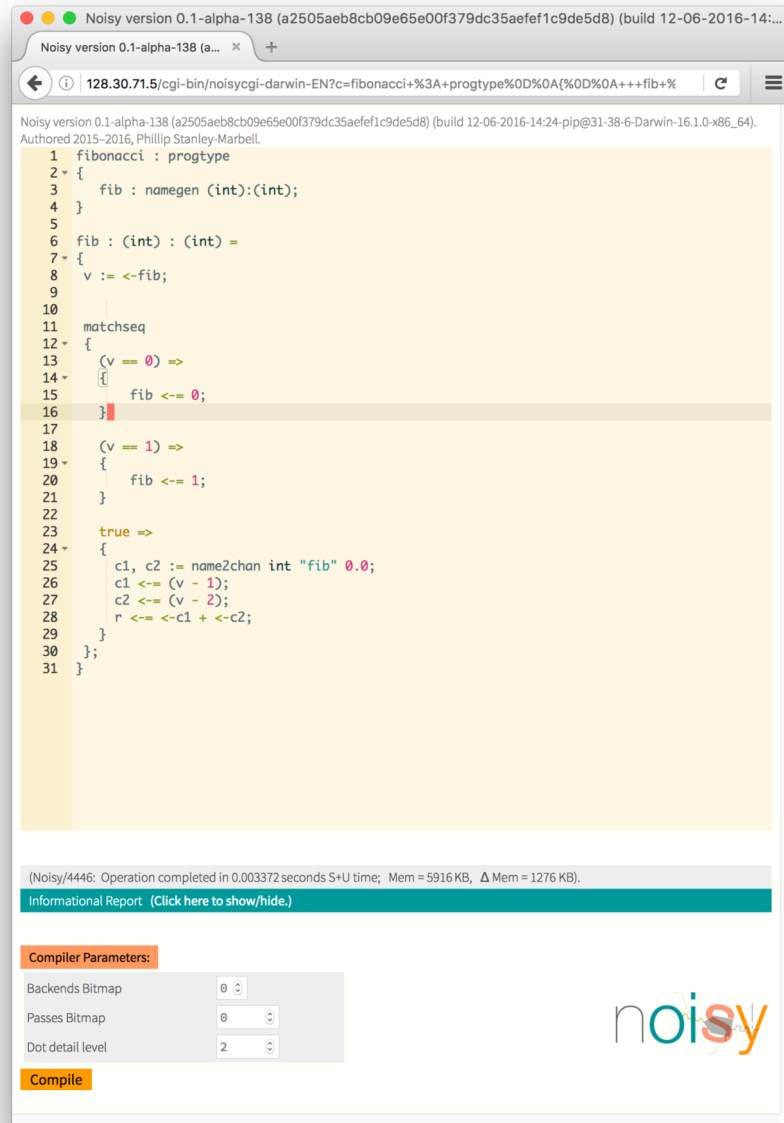


Fig. 22. An online version of the Noisy compiler with a web-based editor allows users to explore writing Noisy programs. The online interface currently only provides a backend to view the program's AST, and does not provide a mechanism to run programs.


```
1 HelloWorld : progtype
2 {
3     init    : namegen ():();
4 }
5
6 init =
7 {
8     print := name2chan System->print "system.print" 0.0;
9     print <-= "Hello World!";
10 }
```

Fig. 23. An example program.



(a) Symbol table state after parsing the HelloWorld program. The Figure is generated automatically by the compiler.

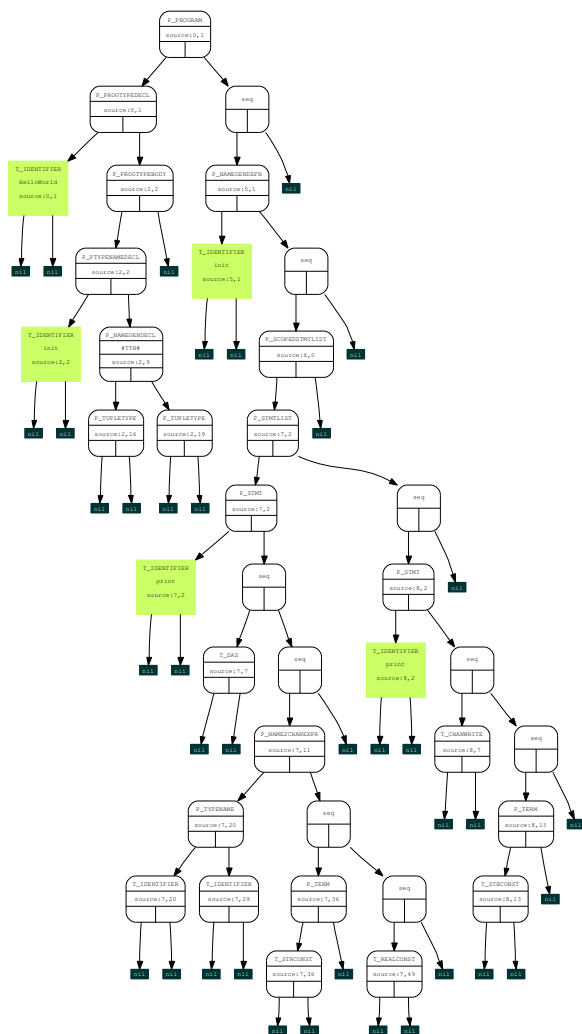


Fig. 24. Intermediate representation (Abstract Syntax Tree (AST)) of HelloWorld program. The Figure is generated automatically by the compiler.